Revision letter
"Expressive power of
linear algebra query languages"

Dear reviewers,

We wish to thank you for your thorough, insightful and constructive comments.

Please find attached our revised manuscript entitled "Expressive power of linear algebra query languages". In the process of revising the paper, we have addressed all your comments and suggestions. Please note that this resulted in several changes to the paper, which we have highlighted in red.

> **Cristian:** For highlight an edit in the document, use the command `EDIT`.

The most important changes are the following.

1. **A clearer write-up of bla bla bla . . .**

   We have completely rewritten bla bla bla . . .

2. **A clearer write-up of bla bla bla . . .**

   We have completely rewritten bla bla bla . . .

Below we also provide a detailed list of answers and explanations to other reviewer comments, excluding those comments that we have already addressed above.

We sincerely hope that this new version of the manuscript sufficiently improves on the previous version, and that it will be considered for publication in TODS.

Best Regards,
The authors.

Reviewer 1

> (. . .) I would however suggest to expand a bit the paragraph on **Related work** at the bottom of page three, and to give more details; in particular this concerns the remark that classical logics with aggregation and fixed-point logic with counting (FPC) can also be used for linear algebra. This point deserves a bit more attention and more references to the literature (for instance to the work of Anuj Dawar and his co-authors). It is known that FPC can express a large collection of linear-algebraic algorithms over fields of characteristic 0 whereas there are severe obstacles in fields with finite characteristics. Incidentally, it is a pity that the authors exclusively consider linear-algebraic problems over the field of reals. I would welcome if the authors could, at least briefly, discuss the power of for-MATLANG for other fields.

ANSWER:**TODO**

Reviewer 2

> (. . .) However, being TODS about Database *systems* and thus usually more oriented towards applied results, I wonder whether an additional (even preliminary) experimental evaluation/implementation of the language would have been needed (I will not fight to have that, though).

ANSWER: We fully agree with the reviewer that an experimental would be useful to have. However, we remark that this paper is a first step in understanding what can be expressed with matrix query languages, and tries to formalize the main concepts, and place them in the context of the existing literature. The paper was invited to TODS as such, and it already has more than 50 pages. We feel that experimental evaluation in fact warrants a paper of its own, and hope to follow up on this in future work. We expanded the "Conclusions and Future Work" section to better reflect this point.

> On the negative side, I believe that the paper is trying, in different places to oversell a bit its results. In particular, the paper makes a big deal in explaining that families of circuits are the de-facto representative logic of linear algebra, but you do not clarify if this is the case under the assumption that, e.g., the depth is bounded. Since your equivalence results are about families of bounded depth/degree, you should properly discuss to what extent, families of circuits of this kind are able to express linear algebra constructs.

ANSWER:**TODO**

> **Cristian:** To what I remember: We agree on lowering the tone with respect to the argument that arithmetic circuits are the de-facto for representing linear algebra. Also, we agree on adding some reference related to writing the determinant with arithmetic circuits (among other operators in linear algebra).

Moreover, although you show via simple examples that there are for-MATLANG expressions (without bound on degree) that are not expressible via families of circuits of polynomial degree, it is not clear whether for-MATLANG (without any restriction) is able to capture all families of circuites without bound on the depth. From what I can see you do not have any proof on whether a family of circuits cannot be expressed via a for-MATLANG expression. This issue should be discussed, either via a theorem, or stating that this problem remains open.

You must be much more transparent in what you are achieving. In different places you claim that you connect for-MATLANG expressions to families of circuits, but then do it in restricted settings. I believe you should make very clear that a full characterization of the form: for any function f over matrices, f is computable by a uniform family of circuits iff it is computable by a for-MATLANG expression is not obtainable, or it is difficult to obtain, and left for future work. Then, you can justify restricting on circuits of bounded depth/degree. Then, the paper should provide a discussion on what these kinds of families can actually express.

This kind of dicussion is particularly relevent in the introduction. The authors should properly clarify what is the actual capacity of for-MATLANG expressions. That is, given the characterization via polynomial degree circuits, which features do you keep, and which are you missing? For example, can you still implement Strassen's algorithm or compute discrete Fourier transformations, as argued in the introduction? I feel that giving the circuit characterization without making explicit what these circuit families can actually do might leave the reader without meaningful "take-home messages", which I guess is what the goal of this paper is: provide key insights on what for-MATLANG can do in terms of *linear algebra* constructs (the equivalence via circuits is "just" the technical tool to convey these messages).

(…)

So, I would request the authors to expand on what for-MATLANG expressions of polynomial degree can atually achieve in terms of lin-

ear algebra constructs (e.g., exploiting the connection with families of bounded degree), or at least state what you cannot achieve. Moreover, it is important to make clear as soon as possible that a connection between general for-MATLANG and general families of circuits is not achievable (e.g., via some some formal statements, or just by making clear ths connection is left as an open problem).

ANSWER:**TODO**

> **Cristian:** With a clear statement that we left for future work the characterization of the expressive power for unrestricted for-Matlang should be enough.

> (. . .) I believe the title is too general. The authors study a *specific* query language, i.e., MATLANG. I understand you study fragments of it, and thus you have languageS, but I feel the title is a bit deceiving. I would make more explicit the content of the paper, specifying it is about the expressive power of MATLANG with iteration.

ANSWER: We agree with this assessment and have changed the title to "Expressive power of MATLANG with bounded iteration".

> When you introduce the $\min(v)$ expression for the first time (after Proposition 3.4) I would anticipate you will explain how to express it in for-MATLANG in the next section.

ANSWER: Done.

> **Thomas:** Moved the clarifying footnote to just before min(v) is used.

> In page 9, definition of $succ(b_i^n, b_j^n)$, I guess you mean $[[succ(u, v)]](I)$, where I maps u and v to b_i^n, and b_j^n. Similarly for $Prev \cdot b_i^n$.

ANSWER: Done.

> Proposition 4.3: here you use the expression "when I assigns V to A". In similar claims, like Proposition 4.2, you do not say anything about what I does to V, and in Proposition 4.1 you use the function mat to state what is the value of V. Please make these equivalence statements more uniform.

ANSWER: Done.

In different parts of the paper you say "circuits of bounded degree". In my view, this usually means that there exists a *constant* that bounds the degree of all circuits in the family, but it is not what you are considering here.

ANSWER: Done

> **Thomas:** Changed to polynomial degree.

Line 22 of Algorithm 1: in the comment I guess meant that $getinput(g)$ outputs i, and not $A[i]$.

ANSWER: Done.

Proposition 5.2 is very long, as it is defining notation in place. I would defined the required notation, such as $vec()$, first, and then give the claim. Moreover, you use Σ, which has never been defined. Do you mean $\{0, 1\}$?

ANSWER: Done.

Proposition 6.3: wherever you use $S(e)$, I guess you mean $type(e)$.

ANSWER: Done.

Figure 4 is again somehow misleading, as you do not prove equivalence with those formalisms in general, but you assume e.g., bounded depth/polynomial degree, binary relations, etc. Please introduce proper notation for these restricted fragments. You could explain this notation in the caption of the figure.

ANSWER:**TODO**

1 Reviewer 3

(...) However, the presentation can still be improved by a lot as indicated below. Also there is some problem in the way how the expressiveness of MATLANG is compared with arithmetic circuits.

(...) While the small errors are easily fixable, the presentation and

the comparison to circuit families of polynomial degree needs a major revision.

ANSWER:**TODO**

> The authors define a MATLANG expression of polynomial degree as any MATLANG expression that has an equivalent circuit family of polynomial degree. Afterwards there is the mind blowing result that this class exactly corresponds to the class of circuits of polynomial degree. Of course this result does not provide any scientific value, as it just repeats the definition. It is not clear at all how this class of MATLANG expressions looks like. Actually it is undecidable if a given MATLANG expression has a polynomial degree.
>
> Instead there should be an (ideally syntactic) definition that is intrinsic to MATLANG. Especially it should not be necessary to refer to circuit families in order to provide a definition of polynomial degree for MATLANG expressions. If no syntactic definition is possible than a sensible semantic definition will also work.

ANSWER:**TODO**

> I start with the comparison to arithmetic circuits. In Section 5.2 you construct MATLANG expressions that uses an input vector of the same arity as the circuit and outputs a single value, the same that the circuit will produce. Theorem 5.1 does not talk about the sizes of other matrices used.
>
> In your construction you use square matrices and vectors of the same size as the input vector. As a result of this design decision, you must limit the construction to circuits of logarithmic depth. For the other direction however, you produce circuits of polynomial depth.
>
> Here you introduce MATLANG expressions of polynomial degree in order to have a MATLANG class and a circuit class of equal expressiveness. Instead, I propose to change your construction of MATLANG expressions in a way that allows to handle all polynomial arithmetic circuits.
>
> Replace Algorithm 1 with an algorithm that computes (and stores) the output values of all gates in topological order. This algorithm is way simpler and does not need a stack.
>
> Then allow your MATLANG construction to use intermediate values of polynomial size. The main data structure is a vector that has as many entries as you have gates in the circuit. Now you can iterate

over all gates (w.l.o.g. the gates are sorted in topological order) and compute all values. This is a single for loop.

Of course you still need the construction from the appendix to compute the nextgate() function, i.e. to simulate the TM that constructs the circuit. Probably it would be simpler if this TM would directly construct a vector that contains all input gates. Then you only need to call this function once for each gate and just use an additional for-loop for the aggregation.

Now you have a natural class of MATLANG expressions that exactly corresponds to arithmetic circuits. I do not see why MATLANG expressions should not be allowed to use intermediate results that have bigger arity than the input. This is a restriction that you never formulated and that is also not imposed on the circuits.

Of course, you can still discuss restricted settings, but please use sensible definitions. E.g. if you restrict the size of intermediate results in MATLANG a corresponding restriction would be on the width of the arithmetic circuit.

ANSWER:**TODO**

You should introduce and describe a consistent notation that allows to easily distinguish whether some MATLANG expressions (1) construct some vector or matrix; or (2) is a Boolean test (i.e., evaluates to a scalar value 0 or 1).

ANSWER:**TODO**

Also, you should adopt the notation that iterator variables are easily distinguishable from other variables throughout the article. Especially if you give expressions like for $col(V, y)$ in line 572, it would be really helpful to immediately see that y is meant to be a variable that can only take canonical vectors as values.

ANSWER:**TODO**

Actually, I would even prefer a notation where iterator variables (like i, j) range over indices, i.e. take values from 1 to n. Then you can still write b_i if you need the canonical vector, but you could also just write V_{ij} instead of $v^t * V * y$, which is way easier to parse for a human. Obviously, such a notation would be just syntactic sugar.

ANSWER:**TODO**

I also suggest to consider introducing some `if` e_1 `then` e_2 `else` e_3 construction, where e_1 is some expression that evaluates to a scalar 0 or 1. This is used a lot in the article.

ANSWER:**TODO**

You have to rename the *succ* and *succ+* expressions, as they in fact do not test for successors. What you call successor is the less or equal relation and what you call *succ+* is the less than relation. So *succ* could be renamed to *islessorequal* and *succ+* could be named *isless*, which also directly reminds the reader that this expression is a Boolean test.

ANSWER:**TODO**

You seem to mix $\rightarrow$ and $\mapsto$ in function specifications in a random way. Function signatures use $\rightarrow$ ($f\colon A \rightarrow B$, where A and B are domain and image of f), while $\mapsto$ is used for concrete mappings (e.g., $f\colon n \mapsto n^2$). And you should use `\colon` instead of :, because : is treated as a division operator by Latex and thus the spacing is not correct.

ANSWER:**TODO**

LU-Decomposition: You should definitely provide some pseudocode for the LU Decomposition algorithm in order to allow a simpler comparison with your MATLANG expressions. Right now the algorithm is given as prose. Furthermore it is not even complete as the definition of c_i with $i \neq 1$ is missing.

ANSWER:**TODO**

Algorithm 1: the aggregate function is working in a completely different way than your MATLANG construction. The MATLANG constructions is a sum over 5 expressions, which especially implies that the order of evaluation is irrelevant. However the algorithm is written in a way that the order of the statements is very important. Especially it is not the case that the five expressions correspond to five different cases of the algorithm as you claim. (. . .) some constructions are way more complex than needed.

ANSWER:**TODO**

Algorithm 1: simplify $e_{iterate}$ by using the loop init $X_1 = e_{start}$. Then you can remove the outer "if-then-else".

ANSWER:**TODO**

Algorithm 1: e_{pop} can just pop both stacks simultaneously by just removing one row from the matrix. The complicated $IdUpTo$ expression is not needed. Just compute

$$e_{pop} := -G_{top} * G_{top}^T * X_k + e_{Prev} * V_{top} * e_{V_{top}}^T$$

to pop both stacks. The first summand computes the delta to remove a line from the matrix and the second summand reads the pointer for the value stack. The pointer for the gate stack will be added in $e_{agg\_(not)\_last}$. Note that this is just a delta. Thus $e_{aggregate}$ will need an additional X_k' summand.

ANSWER:**TODO**

Algorithm 1: e_{agg_prod} also looks way to complicated. Why do you need $IdVal$? Just manipulate the single matrix entry directly:

$$e_{agg\_prod} := V^T * V_{top} * (V^T * (e_{Prev} * V_{top}) - 1)$$

The -1 is to accomodate for the value that is already on the stack. No for-loop needed (hidden in the $IdVal$ expression).

ANSWER:**TODO**

Algorithm 1: I find it confusing that in Figure 3 you list the combined effect of two of the expressions. The figure should describe each of the expressions on its own. Especially as these combinations that you describe are not possible in the algorithm. The algorithm combines always three of the expressions. Also the Figure is wrong. In the upper two cases, there should be no pointer for the gate stack. And in the lower cases there should be an empty value stack according to your construction.

ANSWER:**TODO**

Algorithm 1: Also, please note that you overload n with many different meanings: (1) arity of the circuit, (2) size of the matrices, and (3) length of a bit vector describing a gate id.

ANSWER:**TODO**

> Algorithm 1: Why are gate ids of linear length in the input? This would correspond to an circuit of exponential size.

ANSWER:**TODO**

> Proposition 5.2: First of all, add a runtime bound to the Turing machine. I do not believe that your construction can simulate every linear space machine without restrictions on the runtime.

ANSWER:**TODO**

> Proposition 5.2: And now: Why do you write up the proof using the most complicated type of TMs possible? For your construction it would be perfectly sufficient to have a single tape machine. Your computations take at most two IDs of size $O(log(n))$ as input and produce one such ID as output. This even fits on a single tape if you restrict to length n. You are assuming $n > n_0$ for some sufficiently large n_0 anyway. Using only a single tape makes the construction much more readable, as you can get rid of many indices.

ANSWER:**TODO**

> Proposition 5.2: If you change Algorithm 1 as laid out above, it could make sense to actually allow for an output "tape" that is formed like a square matrix. The head of the output tape could move in four directions (simulated by two vectors wit a single 1 entry). It is obvious (to all that know TMs) that this does not add additional power and one can easily translate the TM producing the circuit to such a TM with a square output. The advantage would be that you could directly produce the adjacency matrix of the circuit which you could use immediately in the expression that evaluates your circuit.

ANSWER:**TODO**

> Proposition 5.2: Also, you do not need to consider the case of small n at all. This case is already considered in the proof of Theorem 1. You can restrict Proposition 5.2 to all n greater than some n_0 (and maybe say that it also holds without this restriction). But no need to waste the space for the proof.

ANSWER:**TODO**

Proposition 5.2: And please try to simplify the construction of the TM. Why do you encode the position of the head in a special way if it is at the edge of the tape? Just adjust the size of the tape such that the end markers are included in the length. Then you do not need to special code this.

ANSWER:**TODO**

From MATLANG to circuits: Why do you restrict the result to MATLANG expressions, where all types only use the size symbol alpha? The construction should work in exactly the same way in the general case. OK, for uniformness you need that all sizes can be computed from the input size by a logspace TM, which results from the definition of uniform circuits where the TM just gets one input parameter. Probably you should discuss this (as it is the usual definition), but in your setting a slightly more general setting of uniform circuits would make sense.

In any case, the proof needs to be reformulated in order to avoid all these pointless case distinctions on the types of the subexpressions. Do all induction cases with $\texttt{type}_S(V) = (\alpha, \beta)$ and provide one induction step for every operation. The cases where one or both of alpha beta are 1 are special cases of the general case and subsumed by the general case. No need to do any case distinctions.

ANSWER:**TODO**

Comparison with K-Relations: Your definition of the renaming operator is nonstandard. Usually this operator takes a function f that renames the variables of a given relation, i.e., the domain of f are the attributes of the relation/expression inside the operator and the image are the new (renamed) attributes. You define it the other way round. When you use the operator, you mix standard and your non-standard definition. Please stay with the established definition.

ANSWER:**TODO**

Comparison with K-Relations: The construction of algebra expression from MATLANG expressions is much more complicated than necessary. You should not do case distinctions on the types of matrices, as the general construction works independently of whether some dimension is 1 or not.

ANSWER:**TODO**

Comparison with K-Relations: Just use $Rel(S)(R_V) := \{row, col\}$ for every matrix V. row and col encode the domain of the indices of the matrix, as in your construction. The only difference is, that this domain could be the singleton $\{1\}$. And you should omit the subscripts alpha and beta of row and col. They are not needed, as the domain is encoded by the relation. If you omit the subscripts you will not need to talk about types at all in most parts of the proof. The soundness of the MATLANG expression ensures that the domains of row and col are correct.

To always provide a col attribute you need a new relation R_1 with attribute col and a single number 1 inside the relation. You can then change $Q(v_p)$ to be

$$\sigma_{row,\gamma_p}(\rho_{\alpha->\gamma_p}(R_\alpha)\texttt{join}\rho_{\alpha->row}(R_\alpha)\texttt{join}R_1)$$

This construction simplifies the definition for transposition to just rename row $\to$ col and col $\to$ row. Also for the other operators you only need to talk about one case. And types are working flawlessly. E.g. matrix product becomes rename col $\to$ C for the first expression and row $\to$ C for the second expression before doing the join. However you have to explain what the join does, i.e., that it just computes the same sum as the matrix product.

ANSWER:**TODO**

line 183: this is ugly to read. Please align at := and at if.

ANSWER:**TODO**

line 216: this also should be aligned.

ANSWER:**TODO**

line 1386: Please rephrase. One could get the impression that the halting problem for linear time TMs is undecidable (which it certainly is not).

ANSWER:**TODO**

line 1435: please provide a pageref for figure 4 or mention that it is at the very end of the article.

ANSWER:**TODO**

line 1443: You have to restrict the expression e, such that it cannot use X. Otherwise Proposition 6.1 is definitely not true.

ANSWER:**TODO**

line 1449: Using your ill-defined definition of expressions of polynomial degree, the proof of Proposition 6.1 is nontrivial and cannot be omitted. You have to show that every expression of sum-MATLANG can be converted to a circuit family of polynomial degree, meeting the syntactic definition of the circuits. Showing that you cannot produce superpolynomial matrix entries is not enough.

ANSWER:**TODO**

line 1681: change Q_1 to Q_2.

ANSWER:**TODO**

line 1731: again, X should not be used inside e.

ANSWER:**TODO**

line 1912 (figure 4): If you adapt the constructions as indicated above, the figure is ok. Otherwise you need to specify which subclass of MATLANG is equivalent to which subclass of arithmetic circuits. E.g., right now, you only can convert MATLANG expressions that use a single size symbol alpha, as you needlessly restrict your construction.

ANSWER:**TODO**

References